

VLSI VAX
Micro-architecture

May 1988

Bob Supnik

For Internal Use Only
Semiconductor Engineering Group

Contents

- The macro-architectural issues
- The canonical micro-architectural model
- MicroVAX
- Improving performance
- CVAX and Rigel
- Summary

The VAX Architecture

- The VAX architecture is a complex instruction set computer (CISC) characterized by:
 - Irregular instruction format (1 to 50+ bytes)
 - Large instruction set (304 instructions)
 - Multiple addressing modes (21)
 - Demand paged virtual memory management
 - Few hardware limitations on software
- TTL/ECL implementations have typically been characterized by:
 - Complex microcode-based control
 - Large control store (400k bits to 2200k bits)
 - Redundant facilities (microcoded and hardware floating point)
 - Inbuilt console I/O
 - Complex memory subsystem (large TB, large cache)
- These implementation characteristics pose severe problems for a single chip VLSI implementation.

The MicroVAX Subset

- Generically, the MicroVAX subset is a set of hardware/ microcode/ software/ performance tradeoffs intended to facilitate VLSI implementation.
- Firmware to software tradeoffs:
 - 59 instructions implemented in macrocode rather than microcode: character string, decimal, EDITPC, CRC, octa-word, h-floating
 - Console implemented in macrocode rather than microcode
- Firmware to hardware tradeoffs:
 - Hardware floating point only
- Performance tradeoffs:
 - Small translation buffer, fully associative, fast replacement
 - No cache or small cache

The Canonical VAX Micro-model

- Most VAX implementations, including the 78X, 750, 730, V-11, MicroVAX, CVAX, Rigel, and Nautilus, have the same basic block structure:
 - I(nstruction) box
 - E(xecution) box
 - M(emory) box
 - Microsequencer/ control store
 - Bus interface unit
 - Interrupts
 - Memory subsystem
 - Console subsystem

I Box

- Parses and decodes instruction stream using internal state and prefetch queue data fetched by the M Box and BIU.
- Gives “VAXness” to the rest of the chip by directing the Microsequencer through specifier evaluation and instruction execution.
- Supplies parameters to specifier evaluation.
- Formats I-stream data for specifier evaluation and instruction execution.

E Box

- Contains main execution data path:
 - Register file
 - ALU and shifter
- Under microcode control, performs:
 - Specifier evaluation
 - Instruction execution
 - Interrupts and exceptions
 - Memory management processing
- Maintains PC, backup PC, PSL, GPRs, RLOG, and other architecturally specified state.

M Box

- Performs address translation and access checking.
- Decodes and initiates memory references, TB accesses.
- Maintains address registers.
- Performs instruction prefetching when idle.

Microsequencer/ Control Store

- Forms next micro-word address and performs micro-word sequencing and access.
- Decodes and selects micro-branch conditions.
- Evaluates requests and initiates micro-traps.
- Maintains micro-stack and pointer.

BIU

- Controls DAL and other external interfaces pins.
- Controls DAL latches and rotators for proper positioning and formatting of incoming and outgoing data.
- Cooperates with M Box in processing of unaligned data.
- Provides autonomous operation on selected I/O functions.

Other

- Interrupt section responds to external hardware and internal software interrupt requests.
- Memory subsystem provides connection of processor to external storage (memory and I/O).
- Console subsystem provides diagnostic and control interface to entire system.
- Note: Console subsystem use **external** in VLSI VAXen and will not be discussed.

Canonical VAX Problems

- All VAX implementations must wrestle with thorny implementation problems posed by the architecture, including:
 - Variable length instructions
 - Unaligned data
 - Virtual memory management
 - Instructions with multiple destinations
 - Instructions with complex algorithms
 - Exceptions
 - Clocking and stalls
- It is interesting to note that there is no reasonable relationship between the difficulty of implementing a feature and its importance.

MicroVAX Overview

- MicroVAX was the first single chip implementation of the VAX. Its characteristics included:
 - Single chip MicroVAX subset CPU (175 instructions) plus companion floating point unit (70 instructions)
 - ZMOS process (3u drawn, NMOS, double level metal)
 - 125,000 transistors, 353 mils x 358 mils
 - 200ns microcycle, 400ns I/O cycle
 - 8 entry TB, no cache
 - 1600 x 39b control store
 - PG - 2/84, LR - 3/85, FRS - 5/85
- MicroVAX implemented a simple external interface:
 - Multiplexed data and address bus (DAL)
 - Address and data strobes (AS, DS)
 - Byte masks for masked writes (BM<3:0>)
 - Cycle status for I/O differentiation (CS<2:0>, WR)
 - DMA request and grant (DMR, DMG)
- MicroVAX drew heavily on the only VLSI full VAX (V-11) for its micro-architectural inspiration (E Box, micro-word, clocking).

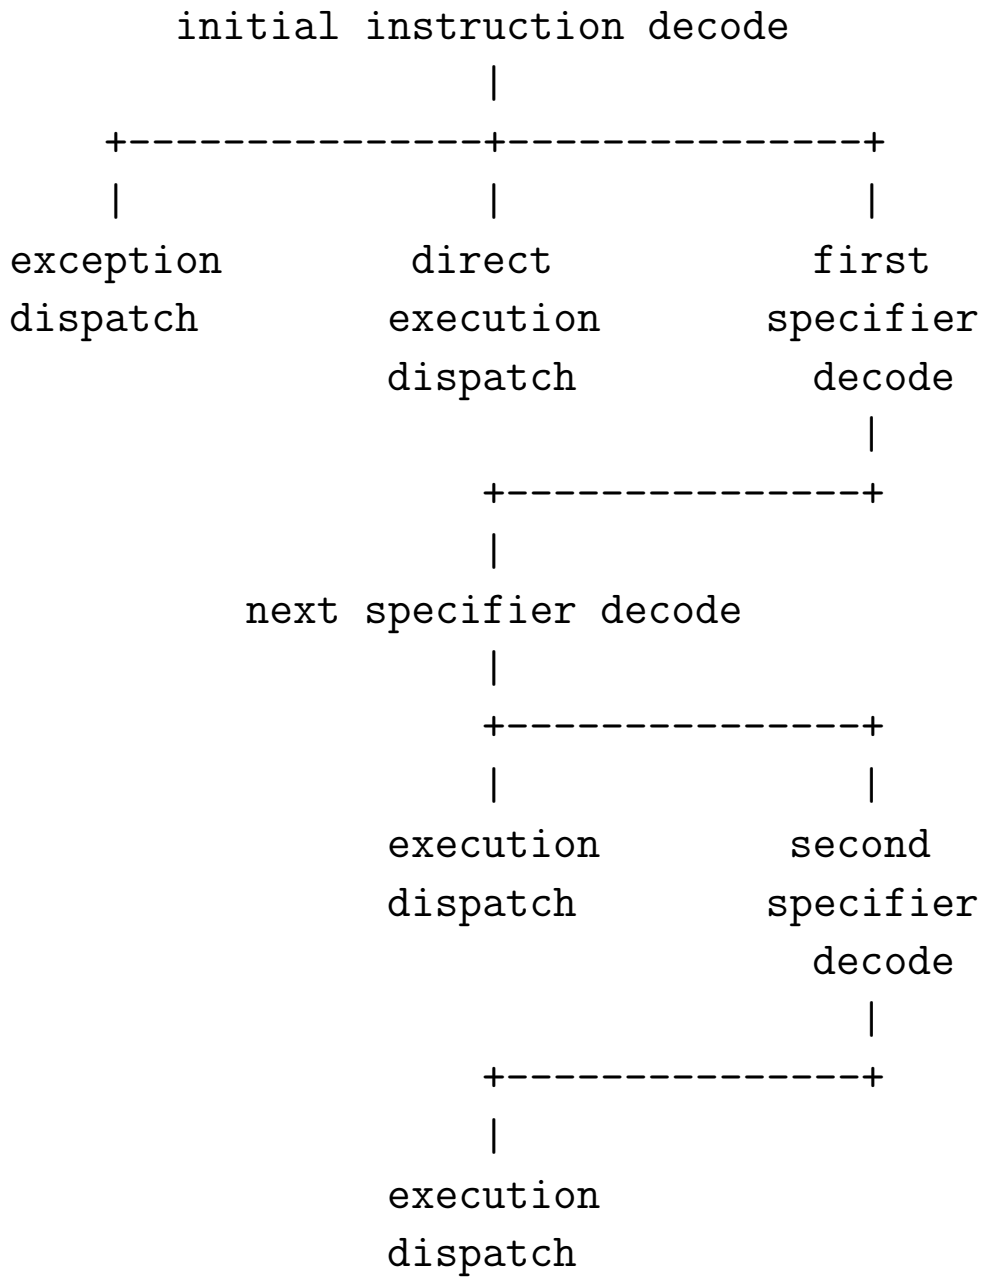
VLSI VAX

Micro-architecture

MicroVAX I Box

- Principal function: prefetch, parse, and decode instructions.
- Prefetch queue:
 - 8 bytes (2 aligned longwords)
 - Maximum of 4 bytes retired per microcycle
- Instruction data register:
 - Data link from I Box to E Box
 - Automatically loaded for simple conditional branches, byte/word displacements, short literals
 - Manually loaded for complex conditional branches, longword displacements, immediates, absolute addresses
- Decode logic:
 - Decode PLAs for exceptions, instructions, specifiers
 - Initial decode PLA (IPLA) to supply instruction parameters to specifier flows
- Control registers:
 - Opcode register
 - Access type/ data length register
 - Current GPR register

MicroVAX Instruction Decode



MicroVAX I Box Dispatches

- Exception dispatches:
 - VAX trap (divide by zero, subscript)
 - Interrupt
 - Trace trap
 - Prefetch exception (no data available)
 - Note: Trap, interrupt dispatch inhibit T bit update
 - Note: Initial instruction decode (IID) can only happen **once**
- Direct execution dispatches:
 - Instructions with no specifiers and simple branches
 - FD prefix
 - PSL<fpd> set
 - Note: FD and FPD interaction

I Box Dispatches, continued

- Specifier dispatches:
 - Short literal
 - Register
 - Indexed
 - Register deferred
 - Autodecrement
 - Autoincrement
 - Immediate
 - Autoincrement deferred
 - Absolute
 - Byte/ word displacement and relative
 - Byte/ word displacement and relative deferred
 - Longword displacement and relative
 - Longword displacement and relative deferred
 - Note: Separate longword dispatch due to 4 bytes per cycle limit on prefetchq queue

MicroVAX Decode Flow

EXAMPLE: ADDL3 R4, 1[^]disp(R5), (R6)+

initial instruction decode --> first spec decode

W[0] := GPR[Rn] ! Flow for first spec

next specifier decode --> second specifier decode

IDR := IB.LONG and case ! Flow for second spec

VA := GPR[Rn] + IDR

W[2] := MEM(VA)

next specifier decode --> execution dispatch

W[0] := W[0] + W[2] ! Execution

microcode dispatch to write destination

VA := GPR[Rn] ! Flow for third spec

GPR[Rn] := GPR[Rn] + 4

MEM(VA) := W[0] ! End of instruction

MicroVAX E Box

- Principal function: Execute VAX macro-instructions.
- Register file:
 - 15 single ported general purpose registers (GPR's): R0 - R14
 - 12 single ported temporary registers (T's): IS, P0BR, P1BR, SBR, SISR, PSL, etc.
 - 7 dual ported working registers (W's): microcode temporaries
- Program counter:
 - Architectural PC (R15)
 - Backup PC, loaded at IID from PC, for exception recovery
 - PC adder, for incrementing PC during instruction parse
- Constant generator:
 - Literal constants from micro-word
 - Fixed constants (0, 1, 4)
 - State dependent constants (KDL, SEXTN)

MicroVAX E Box, continued

- SC/Q register:
 - Working register (W7)
 - Shift counter to control barrel shifter
 - Multiplier/quotient register
 - Case generation register
- Arithmetic/ logical unit (ALU):
 - 32b arithmetic and logic function unit
 - Condition code outputs for 8b, 16b, 32b results
- Barrel shifter:
 - 64b in, 32b out funnel shifter
 - Right shift in hardware
 - Left shift by '32-n' right shift

MicroVAX E Box, continued

- Condition code logic:
 - ‘Raw’ ALU condition codes for microcode testing
 - Architecturally defined PSL condition codes
 - Instruction specified condition code ‘map’ raw condition codes to architecturally specified condition codes
 - Override for developing multi-word condition codes
- Conditional branch logic:
 - Maps opcode against ALU/ PSL condition codes
 - Generates ‘branch taken’ for conditional update of PC
- State logic:
 - Microcode settable/ testable flags
 - Half are global, half are cleared at IID
- Register logging stack:
 - Records autoincrement/ autodecrement modifications to GPRs
 - Used in exception recovery
 - Cleared at IID

MicroVAX M Box

- Principal function: Address translation and external I/O.
- Address registers:
 - VA - address register for data
 - VA' - backup address register for data, autoincrements
 - VIBA - address register for instructions, autoincrements
- Length check logic:
 - SLR, P0LR, P1LR - architecturally specified length registers
 - Length comparator
 - Status output - only tested on TB miss
- Translation buffer:
 - Tag store (CAM) looks up addresses, fully associative
 - Data store (PTEs) holds corresponding PTEs
 - Management algorithm is true LRU

MicroVAX M Box, continued

- Access check logic:
 - Validity check (PTE.V \neq 0) and micro-trap
 - Access (privilege) check and micro-trap
 - M = 0 check and micro-trap
 - Note: probe vs memory request
 - Note: read vs read check, write vs write check
- Unaligned logic:
 - Checks for data transfer across longword boundary
 - Breaks transfer into two transfers with proper data rotation and latching
- Cross page logic:
 - Checks for data transfer across page boundary
 - Initiates micro-trap for proper access check
- Data length logic - drives data length on DAL.
- Micro-trap and abort logic.

Control Store

- Principal function: control memory for chip.
- 1600 words x 39b control store.
- 25b of data path control, in **nine** formats:
 - Basic (ALU)
 - Shift (shifter)
 - Constant (ALU + microcode constant)
 - Special (state twiddling)
 - Mem req (external I/O)
 - MXPR (internal I/O)
 - F Box transfer (FPU I/O)
 - F Box execute (not used)
 - Spare (integer multiply/divide)

Too many! Decoding is a nightmare.

- 14b of sequencing control, in two formats:
 - Jump
 - Branch (conditional or case)

Microsequencer

- Principal function: sequence access of micro-words from control store.
- Provides multiple access modes:
 - Absolute next address
 - Relative next address (signed offset)
 - Sequential next address (micro-PC + 1)
 - Conditional branch
 - Case branch
 - Micro-subroutine and return
 - Externally generated address (test mode)
- Maintains micro-PC (11b), micro-stack (8 entries).
- Mediates and generates micro-traps:
 - M Box - TB miss, ACV/TNV, M = 0, cross page
 - E Box - integer overflow
 - I Box - reserved opcode
 - BIU - floating point error, DAL error

BIU

- Principal function: control external I/O.
- Sequences external I/O functions:
 - Data and interrupt vector read, instruction prefetch
 - Data write with overlap (write and run)
 - FPU transfer
 - DMA request and grant
- Controls data formatting:
 - Write data rotators
 - Read data rotators, latches, zero extender
 - Byte mask pins

Interrupts and Clocks

- Interrupt logic mediates external and internal interrupts:
 - External hardwired interrupts - HALT, PWRFL
 - External vectored interrupts - $IRQ\langle 3:0 \rangle = IPL\langle 17:14 \rangle$
 - Interval timer interrupt and disable flag - $ICCS\langle 6 \rangle$

Internal software interrupts - $SISR\langle 15:1 \rangle = IPL\langle 0F:01 \rangle$ - are implemented entirely in **microcode**.

- Clock logic provides master clocks for all chip logic:
 - Divide by two logic for internal master clock
 - Clock generators for 8 two phase internal clocks
 - Reset and synchronization logic

Too complicated!

Improving Performance

- MicroVAX, like the 11/780, runs at about 500,000 VAX instructions per second:
 - Average 10 microcycles per macro-instruction
 - 200ns microcycles
 - Average macro-instructions is 2.0 microseconds
- To improve performance, there are two, and only two, techniques that can be tried:
 - Shorten the microcycle
 - * by improving technology
 - * by pipelining micro-instructions
 - Reduce the number of microcycles (ticks) per instruction (tpi)
 - * by improved macro-level parallelism
 - * by piecemeal improvement

Faster Microcycles: Technology

- There are four critical loops in a VAX implementation:
 - The E Box loop (register read, ALU, register write)
 - The I Box loop (data in, decode, micro-address out)
 - The Microsequencer loop (control store access, next address decode)
 - The TB/cache loop (address out, translation, access, data in)
- In MicroVAX, each of these loops is balanced around a 200ns period.
- Each generation of technology provides approximately 30% faster gates.
- Therefore, successive generations of VLSI VAXen can speed up by 30% on technology alone.

Can't we do better than that?


VLSI VAX

Micro-architecture


Faster Microcycles: Pipelining

By pipelining the E Box microcycle, micro-instruction throughput can be dramatically increased, thereby reducing the **apparent** microcycle time.

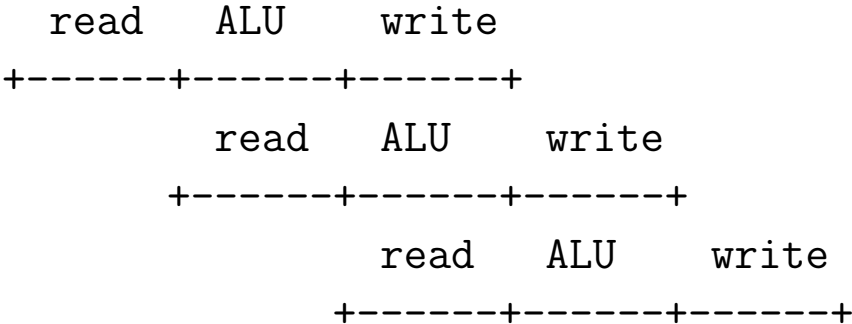
unfolded (1X):



half folded (1.5X):



fully folded (3X):



Micro-pipelining, continued

- Micro-pipelining impacts entire micro-architecture:
 - I Box must be pipelined to meet apparent faster microcycle
 - Microsequencer and control store must get faster to meet apparent faster microcycle
 - TB/cache must get faster to meet apparent faster microcycle
 - Control becomes much more complex throughout due to formal pipeline controls, stalls, etc
 - Microcode becomes much more complex due to longer micro-branch latencies, pipeline side effects, etc
- Micro-pipelining is not a perfect win:
 - Segments are not equal length, effective microcycle time determined by longest segment
 - Pipeline introduces some inefficiencies and stalls
- Micro-pipelining provides the biggest ‘multiplier’ for improving VAX performance; but where do we go after fully folding?

Reduced TPI: Pipelining

- The high TPI of most VAXen is due to two primary factors:
 - Serial decoding of specifiers
 - Lengthy execution times of complex instructions (CALLx, RET, etc)
- Increasing macro-level parallelism could reduce apparent TPI by:
 - Parallel decoding of multiple specifiers, or
 - Overlap of specifier decoding with instruction execution
- **However**, the VAX architecture is highly resistant to macro-level parallelism:
 - Variable length specifiers make parallel decoding of specifiers difficult and expensive
 - Interlocks within and between instructions make overlap of specifiers with instruction execution difficult and expensive
- Most (but not all) VAX architects feel that the costs of macro-level parallelism outweigh the benefits; hence, this approach is not being actively pursued.

Reduced TPI: Nibbling

- If we cannot get a radical reduction in TPI, we can nonetheless get small reductions via piecemeal improvements to the micro-architecture.
- One area for improvement is the memory subsystem. Improvements can include:
 - Enlarged translation buffer
 - On chip cache
 - Multi-level cache
 - Multi-word I/O
 - Write and run (write pipelining)
 - Multiple write buffers
 - Read and run (read pipelining)
 - Hits under misses
- Other areas for improvement:
 - Optimized (via special case) specifier decoding
 - Better hardware support or microcode algorithms for long instructions

CVAX

- Second generation VLSI VAX single chip microprocessor:
 - MicroVAX subset CPU (175 instructions) plus companion floating point unit (70 instructions)
 - CMOS-1 process (2u drawn, CMOS, double level metal)
 - 175,000 transistors, 390 mils x 375 mils
 - 80ns - 100ns microcycle, 160ns - 200ns I/O cycle
 - 28 entry TB, 1kb cache
 - 1600 x 41b control store
- Performance goal is 2.5X - 3.0X current generation:
 - 1.5X from technology improvements
 - 1.5X from micro-architectural pipelining
 - Remainder from improved memory subsystem

VLSI VAX

Micro-architecture

CVAX, continued

- Faster microcycle – technology:
 - CMOS-1 process substantially faster than ZMOS (2ns representative gate delay vs 3ns)
 - Lower power permits fuller use of large devices for speed-critical paths
- Faster microcycle – micro-pipelining:
 - Half folded micro-pipeline
 - Register file writes through, thereby allowing writes under reads with no explicit bypass logic
 - RAS/CAS addressing of control store provides same micro-branch latency as in MicroVAX (one cycle)
 - Pipeline in I Box adds one cycle to macro-branch latency
- Reduced TPI – better memory subsystem:
 - Enlarged TB (28 entries vs 8) for reduced misses
 - On chip single cycle cache (1kb, two way associative, 8 byte block)
 - Off chip two cycle cache (64kb+, direct map)
 - Multi-word read for on chip cache fill

CVAX, continued

- CVAX I Box is based on Nautilus rather than 780:
 - I Box is an autonomous state machine which parses the instruction stream based on its own state data
 - I Box parses **all** specifiers using one generic (parameterized) set of specifier flows
 - I Box and E Box are synchronized by a single directive, **DECODER NEXT**
 - Prefetch queue is 12 bytes (3 aligned longwords), allowing retirement of up to 6 bytes per microcycle
 - Instruction data register automatically loaded in most cases (only immediates and complex branch displacements are done manually)
- CVAX E Box implements half folded micro-pipeline:
 - All registers have extra (write) port
 - Writes are executed under reads, with bypass through the register file
 - 4 extra T registers for per process stack pointers
 - SC and Q are separate registers
 - PSL is maintained in hardware
- CVAX M Box is like MicroVAX:
 - 28 TB entries
 - Not last used (NLU) replacement algorithm

CVAX, continued

- CVAX control store and Microsequencer are simplified:
 - 1600 x 41b control store
 - Five (rather than nine) data path formats
 - Two sequencing formats
 - Paged rather than signed displacement addressing
 - Case rather than conditional branching
 - 8 way rather than 16 way cases
- CVAX BIU provides increased flexibility:
 - On chip 1kb single cycle cache
 - Multi-word reads for cache fills
 - Externally requested cycle retry
 - Optional data parity
 - Much more efficient FPA protocol
- Improvements in interrupts and clocking:
 - Two more hardwired interrupts (CRD, MEMERR)
 - Partial hardware implementation of software interrupts
 - Externally generated four phase overlapping clocks

CVAX Decode Flow

EXAMPLE: ADDL3 R4, 1[^]disp(R5), (R6)+

decoder next --> specifier decode

W[Sn] := GPR[Rn] ! Flow for specifier
! Sn = 0, Rn = 4

decoder next --> specifier decode (IDR loaded)

VA := GPR[Rn] + IDR ! Flow for specifier
W[Sn] := MEM(VA) ! Sn = 2, Rn = 5

decoder next --> specifier decode

VA, W[Sn] := GPR[Rn] ! Flow for specifier
GPR[Rn] := GPR[Rn] + 4 ! Sn = 4, Rn = 6

decoder next --> execution dispatch

W[0] := W[0] + W[2] ! Execution
MEM(VA) := W[0] ! End of instruction

Rigel

- Third generation VLSI VAX single chip microprocessor:
 - MicroVAX subset CPU (175 instructions) plus companion floating point unit (70 instructions)
 - CMOS-2 process (1.5u drawn, CMOS, double level metal)
 - 325,000 transistors, tbd mils x tbd mils
 - 30 ns - 40ns microcycle, 90ns - 120ns I/O cycle
 - 64 entry TB, 2kb cache
 - 1700 x 50b control store
- Performance goal is 6X - 8X current generation:
 - 2X from technology improvements
 - 3X from micro-architectural pipelining
 - Remainder from improved memory subsystem

VLSI VAX

Micro-architecture

Rigel

- Faster microcycle – technology:
 - CMOS-2 process substantially faster than ZMOS (1.5ns representative gate delay vs 3ns)
 - Lower power permits fuller use of large devices for speed-critical paths
- Faster microcycle – micro-pipelining:
 - Fully folded micro-pipeline
 - Register file writes through, thereby allowing writes under reads with just one level of explicit bypass logic
 - Micro-branch latency increases to three cycles
 - Pipeline in I Box adds yet another cycle to macro-branch latency
- Reduced TPI – better memory subsystem:
 - Enlarged TB (64 entries) for reduced misses
 - On chip single cycle cache (2kb, direct map, 8 byte block)
 - Off chip three cycle cache (128kb, direct map, 16 byte fill size, 64 byte block size)
 - Multi-word read for all cache fills
 - Multi-word writes for burst output situations
 - Read and run pipeline

Rigel, continued

- Rigel I Box is based on CVAX/ Nautilus rather than 780:
 - I Box is an autonomous state machine which parses the instruction stream based on its own state data
 - I Box parses **all** specifiers using one generic (parameterized) set of specifier flows
 - I Box and E Box are synchronized by a single directive, **DECODER NEXT**
 - Prefetch queue is 16 bytes (4 aligned longwords), allowing retirement of up to 10 bytes per microcycle
 - Instruction data register automatically loaded in **all** cases
- Rigel E Box implements fully folded micro-pipeline, plus read pipelining:
 - All registers have extra (write) port
 - MD (working) registers have second write port plus valid bits for synchronization
 - Bypass around ALU/ shifter and through register file
 - 8 extra T registers for per process stack pointers and memory management length registers
 - MD7 is separate register, SC and Q are again combined
 - PSL is maintained in hardware
- Rigel M Box is simplified:
 - 64 TB entries

- Not last used (NLU) replacement algorithm
- Length checks implemented in microcode rather than in hardware

Rigel, continued

- Rigel control store and Microsequencer are simplified:
 - 1600 x 50b control store
 - Four data path formats
 - Two sequencing formats
 - Paged rather than signed displacement addressing
 - Case rather than conditional branching
 - 8 way rather than 16 way cases
- BIU provides even more flexibility:
 - On chip 2kb single cycle cache
 - Multi-word reads for cache fills
 - Multi-word writes for high output
 - Externally requested cycle retry
 - Mandatory data parity
 - Much more efficient FPA protocol
- Improvements in interrupts and clocking:
 - Two more hardwired interrupts (CRD, MEMERR)
 - Full hardware implementation of software interrupts
 - Externally generated four phase overlapping clocks

Rigel Decode Flow

EXAMPLE: ADDL3 R4, 1[^]disp(R5), (R6)+

decoder next --> specifier decode

MD[Sn] := GPR[Rn] ! Flow for specifier
! Sn = 0, Rn = 4

decoder next --> specifier decode (IDR loaded)

MD[Sn] := MEM(GPR[Rn]+IDR) ! Flow for specifier
! Sn = 2, Rn = 5

decoder next --> specifier decode

VA := GPR[Rn] ! Flow for specifier
GPR[Rn] := GPR[Rn] + 4 ! Sn = 4, Rn = 6

decoder next --> execution dispatch

MEM(VA) := MD[0] + MD[2] ! Execution

Summary

- The implementation of the VAX in VLSI has required some adaptations and adjustments at the macro-architectural level.
- The four VLSI VAXen defined to date (MicroVAX, V-11, CVAX, and Rigel) all follow the same (canonical) micro-architectural model.
- The implementation process is complex, with much effort expended on architectural nits that have little or no performance benefit.
- The constraints of the VAX architecture have limited attempts at performance improvement to just three basic areas:
 - Improved technology
 - Microcycle pipelining
 - Improved memory subsystem
- Despite the difficulties, the VLSI VAXen have proven both popular and competitive, and will form the basis of DEC's low end and mid range product offerings for years to come.